

Q1. HTSPC- Hate Speech Classification

Packages/libraries needed to be installed :

- Natural Language Toolkit (NLTK) library
- tweet-preprocessor module
- wordninja

Input format :

We use command line arguments to take the input of the paths of the files.

We take the path of the training data csv file as our first command line argument and path of test data file as our second argument.

Data preprocessing :

1. First we use the tweet preprocessor library in Python to remove the URLs, emojis, numbers and mention tags present in the data.

The code snippet is:

▼ In this section we are removing all URLs, emojis, numbers and mention tags from the data.

```
import preprocessor as p
p.set_options(p.OPT.URL, p.OPT.EMOJI, p.OPT.SMILEY, p.OPT.NUMBER, p.OPT.MENTION)
i = 0
for line in X:
    cleaned_line = p.clean(line)
    X[i] = cleaned_line
    i = i+1
```

2. We use the wordninja module to separate the different words.

▼ Using the wordninja module, we are splitting the data into separate words.

```
[46] import wordninja
i = 0
for text in X:
    arr = wordninja.split(text)
    arr1 = ""
    for j in arr:
        arr1 = arr1+" "+j
    X[i] = arr1
    i = i+1
```

3. We remove the special characters and extra spaces from the data by writing regular expressions and matching them to the strings.

- We also test using CountVectorizer.
- We test using different ngram ranges.

Model selection :

Now, we perform train validation data split .

After that, we run different models to test which one perform better.

The outcomes are shown below :

1. Using logistic regression model with different values of C

▼ Logistic Regression model

```
[113] from sklearn.linear_model import LogisticRegression

c = [0.001, 0.01, 0.05, 0.5, 0.1, 1.0]
for i in c:
    model_lr = LogisticRegression(C=i)
    model_lr.fit(X_train, y_train)
    y_pred = model_lr.predict(X_test)
    accuracy_lr = accuracy_score(y_test, y_pred)
    f1_lr = f1_score(y_test, y_pred)
    print ("C = ", i, " Accuracy: ", accuracy_lr)
    print ("C = ", i, " F1 score: ", f1_lr)
```

```
↳ C = 0.001 Accuracy: 0.635673624288425
C = 0.001 F1 score: 0.7735849056603773
C = 0.01 Accuracy: 0.6603415559772297
C = 0.01 F1 score: 0.7716836734693877
C = 0.05 Accuracy: 0.6802656546489564
C = 0.05 F1 score: 0.7715254237288135
C = 0.5 Accuracy: 0.6622390891840607
C = 0.5 F1 score: 0.751048951048951
C = 0.1 Accuracy: 0.6802656546489564
C = 0.1 F1 score: 0.7696514012303487
C = 1.0 Accuracy: 0.6574952561669829
C = 1.0 F1 score: 0.7448763250883393
```

2. Using Linear SVC model with different values of C

Linear SVC model

```
from sklearn.svm import LinearSVC

c = [0.001, 0.01, 0.05, 0.5, 0.1, 1.0]
for i in c:
    model_svm = LinearSVC(C=i)
    model_svm.fit(X_train, y_train)
    y_pred = model_svm.predict(X_test)
    accuracy_svm = accuracy_score(y_test, y_pred)
    f1_svm = f1_score(y_test, y_pred)
    print("C = ", i, " Accuracy: ", accuracy_svm)
    print("C = ", i, " F1 score: ", f1_svm)
```

```
C = 0.001 Accuracy: 0.6660341555977229
C = 0.001 F1 score: 0.7740693196405648
C = 0.01 Accuracy: 0.6755218216318786
C = 0.01 F1 score: 0.7663934426229508
C = 0.05 Accuracy: 0.6574952561669829
C = 0.05 F1 score: 0.7477288609364081
C = 0.5 Accuracy: 0.6404174573055028
C = 0.5 F1 score: 0.7294789436117058
C = 0.1 Accuracy: 0.6584440227703985
C = 0.1 F1 score: 0.7457627118644068
C = 1.0 Accuracy: 0.6299810246679317
C = 1.0 F1 score: 0.7202295552367288
```

3. Using SVM with 'linear' kernel

SVM with 'linear' kernel

```
svm = SVC(kernel = 'linear')
svm.fit(X_train, y_train)
y_pred = svm.predict(X_test)
accuracy_svm = accuracy_score(y_test, y_pred)
f1_svm = f1_score(y_test, y_pred)
print("Accuracy: ", accuracy_svm)
print(" F1 score: ", f1_svm)
```

```
Accuracy: 0.6328273244781784
F1 score: 0.7205776173285199
```

4. Using SVM with 'poly' kernel

SVM with 'poly' kernel

```
[134] svm = SVC(kernel = 'poly')
      svm.fit(X_train, y_train)
      y_pred = svm.predict(X_test)
      accuracy_svm = accuracy_score(y_test, y_pred)
      f1_svm = f1_score(y_test, y_pred)
      print("Accuracy: ", accuracy_svm)
      print(" F1 score: ", f1_svm)
```

```
Accuracy: 0.6385199240986718
F1 score: 0.7700663850331926
```

5. Using SVM with 'rbf' kernel

▾ SVM with 'rbf' kernel

```
[132] svm = SVC(kernel = 'rbf')
      svm.fit(X_train, y_train)
      y_pred = svm.predict(X_test)
      accuracy_svm = accuracy_score(y_test, y_pred)
      f1_svm = f1_score(y_test, y_pred)
      print ("Accuracy: ", accuracy_svm)
      print (" F1 score: ", f1_svm)
```

```
➤ Accuracy:  0.6717267552182163
  F1 score:  0.7747395833333333
```

From the above outcomes we choose SVM with 'rbf' kernel as our final model, as it performs pretty well.

We load the test data, and apply the same preprocessing on it, as we applied on the train data.

Finally, we save our predictions in a csv file.

We get an accuracy of around 80% on the final test data.